




# CLUSTERING



# WHAT IS CLUSTER ANALYSIS?

- Given a set of objects, place them in groups such that:
  - the objects in a group are similar (or related)
  - different from (or unrelated to) the objects in other groups



# APPLICATIONS OF CLUSTER ANALYSIS


## ■ Understanding

- Group related documents for browsing, group genes and proteins that have similar functionality, or group stocks with similar price fluctuations




## ■ Summarization


- Reduce the size of large data sets





# NOTION OF A CLUSTER CAN BE AMBIGUOUS





How many clusters?




Six Clusters



Two Clusters




Four Clusters





# TYPES OF CLUSTERINGS



- A **clustering** is a set of clusters
- Important distinction between **hierarchical** and **partitional** sets of clusters
  - Partitional Clustering
    - A division of data objects into non-overlapping subsets (clusters)
  - Hierarchical clustering
    - A set of nested clusters organized as a hierarchical tree



# PARTITIONAL CLUSTERING




**Original Points**




**A Partitional Clustering**


# HIERARCHICAL CLUSTERING




**Traditional Hierarchical Clustering**



**Traditional Dendrogram**




**Non-traditional Hierarchical Clustering**




**Non-traditional Dendrogram**

# TYPES OF CLUSTERING

Partitional Clustering



Hierarchical Clustering





# OTHER DISTINCTIONS BETWEEN SETS OF CLUSTERS

- Exclusive versus non-exclusive

- non-exclusive clustering:

- points may belong to multiple clusters.

- an belong to multiple classes or could be 'border' points

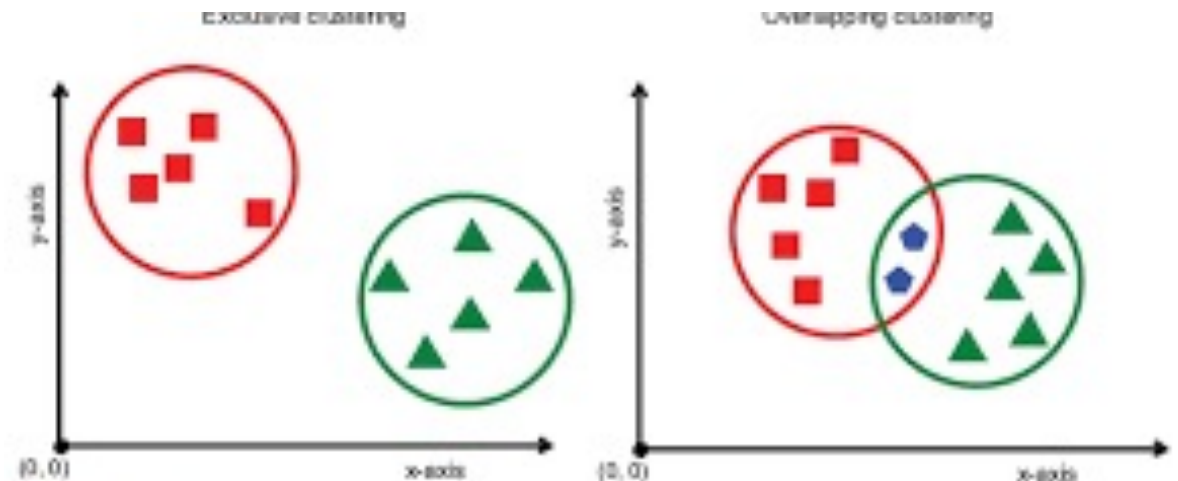
- fuzzy clustering : a point belongs to every cluster with some weight between 0 and 1

- weights sum to 1


- probabilistic clustering has similar characteristics

- Partial versus complete


- Partial: only cluster some of the data




## OTHER DISTINCTIONS BETWEEN SETS OF CLUSTERS




Raw Data



Clustered Data



Raw Data




Clustered Data

# TYPES OF CLUSTERS

- Well-separated clusters
- Prototype-based clusters
- Contiguity-based clusters
- Density-based clusters
- Described by an Objective Function

## TYPES OF CLUSTERS: WELL-SEPARATED

- Well-Separated Clusters:
  - A cluster is a set of points such that any point in a cluster is closer to every other point in the cluster than to any point not in the cluster.




**3 well-separated clusters**

# TYPES OF CLUSTERS: PROTOTYPE-BASED

- Prototype-based





- A cluster is a set of objects such that an object in a cluster is closer (more similar) to the prototype or “center” of a cluster, than to the center of any other cluster
- **centroid**, the average of all the points in the cluster,
- **medoid**, the most “representative” point of a cluster



**4 center-based clusters**

## TYPES OF CLUSTERS: CONTIGUITY-BASED


- Contiguous Cluster (Nearest neighbor or Transitive)
  - A cluster is a set of points such that a point in a cluster is closer to one or more other points in the cluster than to any point not in the cluster.



**8 contiguous clusters**

## TYPES OF CLUSTERS: DENSITY-BASED

- Density-based
  - A cluster is a dense region of points, which is separated by low-density regions, from other regions of high density.
  - Used when the clusters are irregular or intertwined, and when noise and outliers are present.



**6 density-based clusters**

# TYPES OF CLUSTERS: OBJECTIVE FUNCTION

- Clusters Defined by an Objective Function
  - Finds clusters that minimize or maximize an objective function.
  - Enumerate all possible ways of dividing the points into clusters
  - Evaluate the goodness of each potential set of clusters.
  - Objectives can be global or local
    - Hierarchical clustering algorithms typically have local objectives
    - Partitional algorithms typically have global objectives



# CHARACTERISTICS OF THE INPUT DATA ARE IMPORTANT

- Type of proximity or density measure
  - Central to clustering
  - Depends on data and application
- Data characteristics that affect proximity and/or density are
  - Dimensionality
    - Sparseness
  - Attribute type
  - Special relationships in the data
    - For example, autocorrelation
  - Distribution of the data
- Noise and Outliers
  - Often interfere with the operation of the clustering algorithm
- Clusters of differing sizes, densities, and shapes

# CLUSTERING ALGORITHMS


- K-means and its variants
- Hierarchical clustering
- Density-based clustering

# K-MEANS CLUSTERING







- Partitional clustering approach
- Number of clusters,  $K$ , must be specified
- Each cluster is associated with a **centroid**
- Each point is assigned to the cluster with the closest centroid
- The basic algorithm is very simple

- 
- 1: Select  $K$  points as the initial centroids.
  - 2: **repeat**
  - 3:   Form  $K$  clusters by assigning all points to the closest centroid.
  - 4:   Recompute the centroid of each cluster.
  - 5: **until** The centroids don't change
-

# EXAMPLE OF K-MEANS CLUSTERING



# EXAMPLE OF K-MEANS CLUSTERING



# K-MEANS CLUSTERING – DETAILS

- Simple iterative algorithm.
  - Choose initial centroids;
  - repeat {assign each point to a nearest centroid; re-compute cluster centroids}
  - until centroids stop changing.
- Initial centroids are often chosen randomly.
  - Clusters produced can vary from one run to another
- The centroid is (typically) the mean of the points in the cluster, but other definitions are possible
- K-means will converge for common proximity measures with appropriately defined centroid
- Most of the convergence happens in the first few iterations.
  - Often the stopping condition is changed to 'Until relatively few points change clusters'
- Complexity is  $O(n * K * I * d)$ 
  - $n$  = number of points,  $K$  = number of clusters,  
 $I$  = number of iterations,  $d$  = number of attributes


# K-MEANS OBJECTIVE FUNCTION

- A common objective function (used with Euclidean distance measure) is Sum of Squared Error (SSE)
  - For each point, the error is the distance to the nearest cluster center
  - To get SSE, we square these errors and sum them.


$$SSE = \sum_{i=1}^K \sum_{x \in C_i} dist^2(m_i, x)$$

- $x$  is a data point in cluster  $C_i$  and  $m_i$  is the centroid (mean) for cluster  $C_i$
- SSE improves in each iteration of K-means until it reaches a local or global minima.


# TWO DIFFERENT K-MEANS CLUSTERING



**Optimal Clustering**




**Original Points**



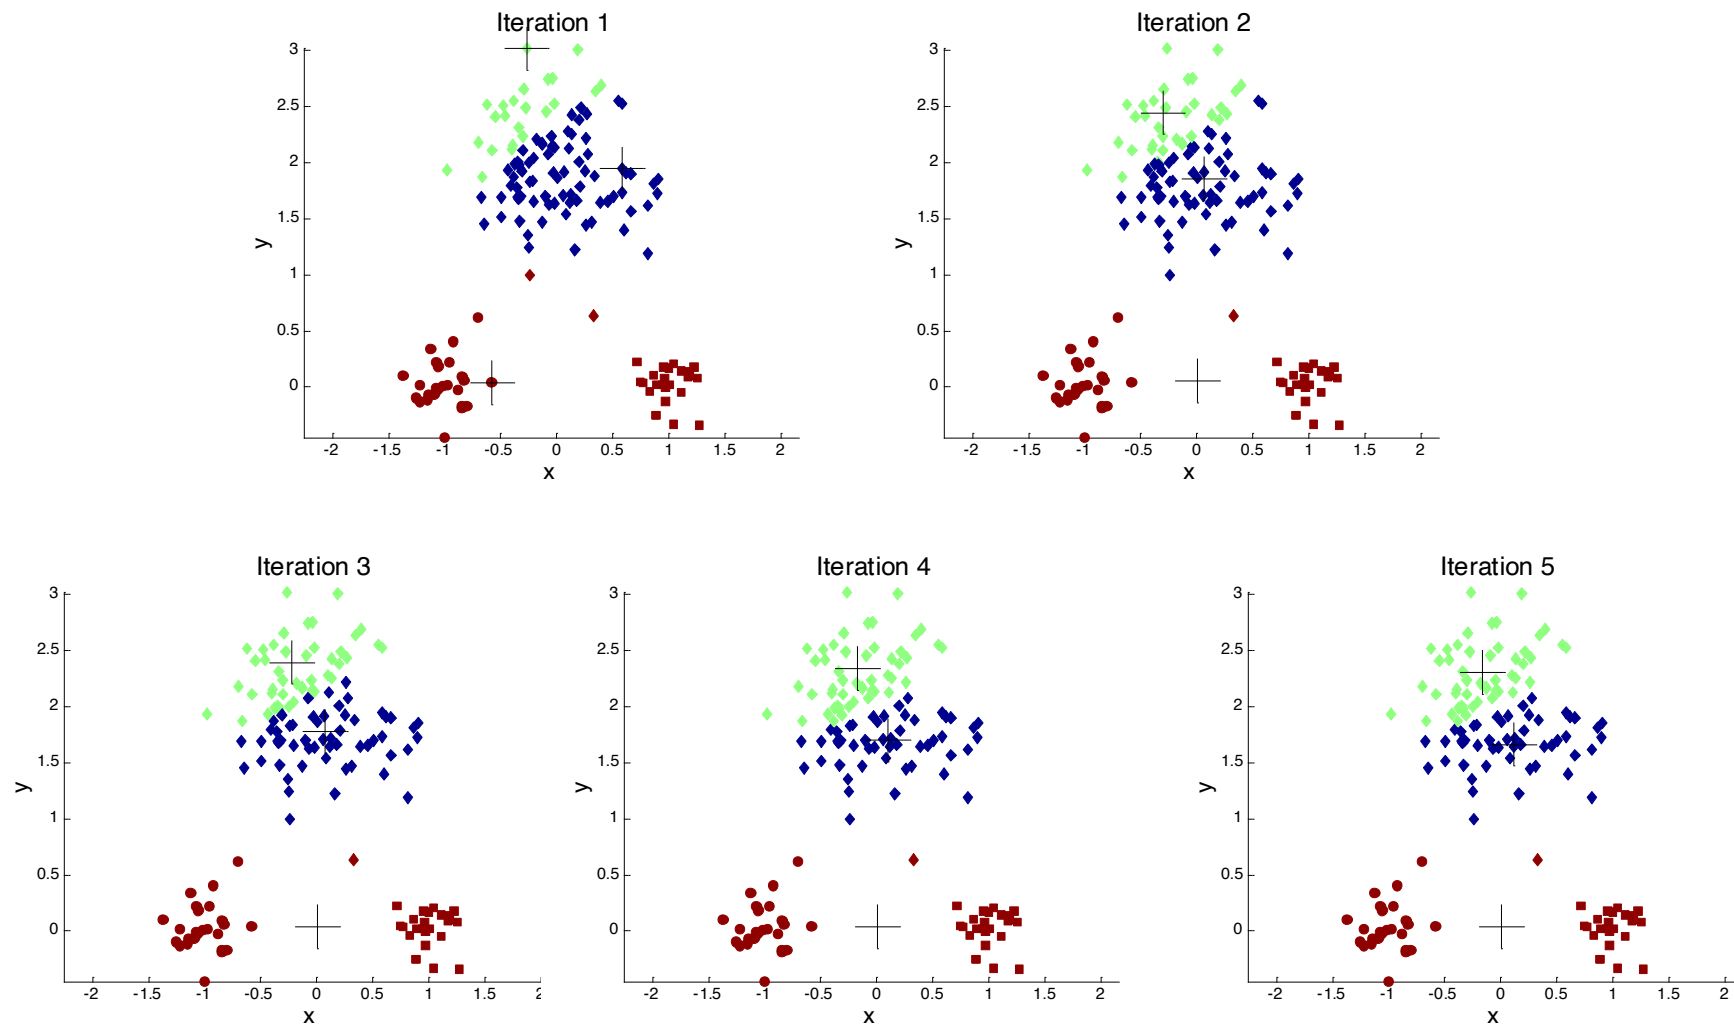
**Sub-optimal Clustering**



## IMPORTANCE OF CHOOSING INITIAL CENTROIDS ...



# IMPORTANCE OF CHOOSING INITIAL CENTROIDS ...




# PROBLEMS WITH SELECTING INITIAL POINTS

- If there are  $K$  'real' clusters, then the chance of selecting one centroid from each cluster is small.
- Chance is relatively small when  $K$  is large
- If clusters are the same size  $n$  then

$$P = \frac{\text{number of ways to select one centroid from each cluster}}{\text{number of ways to select } K \text{ centroids}} = \frac{K!n^K}{(Kn)^K} = \frac{K!}{K^K}$$





- For example, if  $K = 10$ , then probability =  $10!/10^{10} = 0.00036$

# 10 CLUSTERS EXAMPLE




**Starting with two initial centroids in one cluster of each pair of clusters**

# 10 CLUSTERS EXAMPLE







**Starting with two initial centroids in one cluster of each pair of clusters**

# 10 CLUSTERS EXAMPLE



**Starting with some pairs of clusters having three initial centroids, while other have only one.**

# 10 CLUSTERS EXAMPLE



Starting with some pairs of clusters having three initial centroids, while other have only one.

# SOLUTIONS TO INITIAL CENTROIDS PROBLEM

- Multiple runs
- Use some strategy to select the  $k$  initial centroids and then select among these initial centroids
  - Select most widely separated
  - K-means++ is a robust way of doing this selection
  - Use hierarchical clustering to determine initial centroids
- Bisecting K-means
  - Not as susceptible to initialization issues



# K-MEANS++

- The k-means++ algorithm guarantees an approximation ratio  $O(\log k)$  in expectation, where  $k$  is the number of centers

To select a set of initial centroids,  $C$ , perform the following

Select an initial point at random to be the first centroid

For  $k - 1$  steps





For each of the  $N$  points,  $x_i$ ,  $1 \leq i \leq N$ , find the minimum squared distance to the currently selected centroids,

$C_1, \dots, C_j$ ,  $1 \leq j < k$ , i.e.,  $\min_j d^2(C_j, x_i)$

Randomly select a new centroid by choosing a point with probability proportional to  $\frac{\min_j d^2(C_j, x_i)}{\sum_i \min_j d^2(C_j, x_i)}$

End For


# K-MEAN++



# BISECTING K-MEANS

- Bisecting K-means algorithm
  - Variant of K-means that can produce a partitional or a hierarchical clustering
- 
- 1: Initialize the list of clusters to contain the cluster containing all points.
  - 2: **repeat**
  - 3:   Select a cluster from the list of clusters
  - 4:   **for**  $i = 1$  to *number\_of\_iterations* **do**
  - 5:     Bisect the selected cluster using basic K-means
  - 6:   **end for**
  - 7:   Add the two clusters from the bisection with the lowest SSE to the list of clusters.
  - 8: **until** Until the list of clusters contains  $K$  clusters
-


# BISECTING K-MEANS EXAMPLE




# LIMITATIONS OF K-MEANS

- K-means has problems when clusters are of differing
  - Sizes
  - Densities
  - Non-globular shapes
- K-means has problems when the data contains outliers.
  - One possible solution is to remove outliers before clustering

# LIMITATIONS OF K-MEANS: DIFFERING SIZES




**Original Points**




**K-means (3 Clusters)**

# LIMITATIONS OF K-MEANS: DIFFERING DENSITY




**Original Points**




**K-means (3 Clusters)**

# LIMITATIONS OF K-MEANS: NON-GLOBULAR SHAPES




**Original Points**




**K-means (2 Clusters)**



# OVERCOMING K-MEANS LIMITATIONS




**Original Points**




**K-means Clusters**

One solution is to find a large number of clusters such that each of them represents a part of a natural cluster. But these small clusters need to be put together in a post-processing step.

# OVERCOMING K-MEANS LIMITATIONS




**Original Points**




**K-means Clusters**

One solution is to find a large number of clusters such that each of them represents a part of a natural cluster. But these small clusters need to be put together in a post-processing step.

# OVERCOMING K-MEANS LIMITATIONS



**Original Points**



**K-means Clusters**

One solution is to find a large number of clusters such that each of them represents a part of a natural cluster. But these small clusters need to be put together in a post-processing step.