



CLASSIFICATION

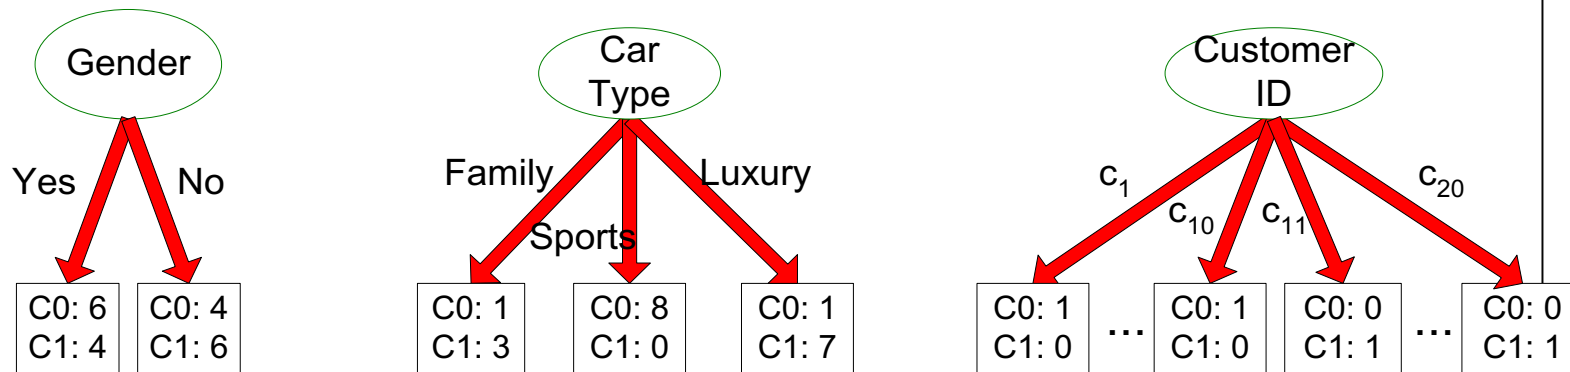


HOW TO DETERMINE THE BEST SPLIT

Before Splitting: 10 records of class 0 (c0),
10 records of class 1 (c1)

What are the values of the label for this data? How many cases / records for each label.

Learn the type of each attribute / feature, their values.



Which test condition is the best?

Customer Id	Gender	Car Type	Shirt Size	Class
1	M	Family	Small	C0
2	M	Sports	Medium	C0
3	M	Sports	Medium	C0
4	M	Sports	Large	C0
5	M	Sports	Extra Large	C0
6	M	Sports	Extra Large	C0
7	F	Sports	Small	C0
8	F	Sports	Small	C0
9	F	Sports	Medium	C0
10	F	Luxury	Large	C0
11	M	Family	Large	C1
12	M	Family	Extra Large	C1
13	M	Family	Medium	C1
14	M	Luxury	Extra Large	C1
15	F	Luxury	Small	C1
16	F	Luxury	Small	C1
17	F	Luxury	Medium	C1
18	F	Luxury	Medium	C1
19	F	Luxury	Medium	C1
20	F	Luxury	Large	C1

MEASURES OF NODE IMPURITY

● Gini Index

$$Gini\ Index = 1 - \sum_{i=0}^{c-1} p_i(t)^2$$

Where $p_i(t)$ is the frequency of class i at node t , and c is the total number of classes

● Entropy

$$Entropy = - \sum_{i=0}^{c-1} p_i(t) \log_2 p_i(t)$$

● Misclassification error (confusing matrix for decision tree)

$$Classification\ error = 1 - \max[p_i(t)]$$

FINDING THE BEST SPLIT


1. Compute impurity measure (P) before splitting
2. Compute impurity measure (M) after splitting
 - Compute impurity measure of each child node
 - M is the weighted impurity of child nodes
3. Choose the attribute test condition that produces the highest gain

$$\text{Gain} = P - M$$

or equivalently, lowest impurity measure after splitting (M)

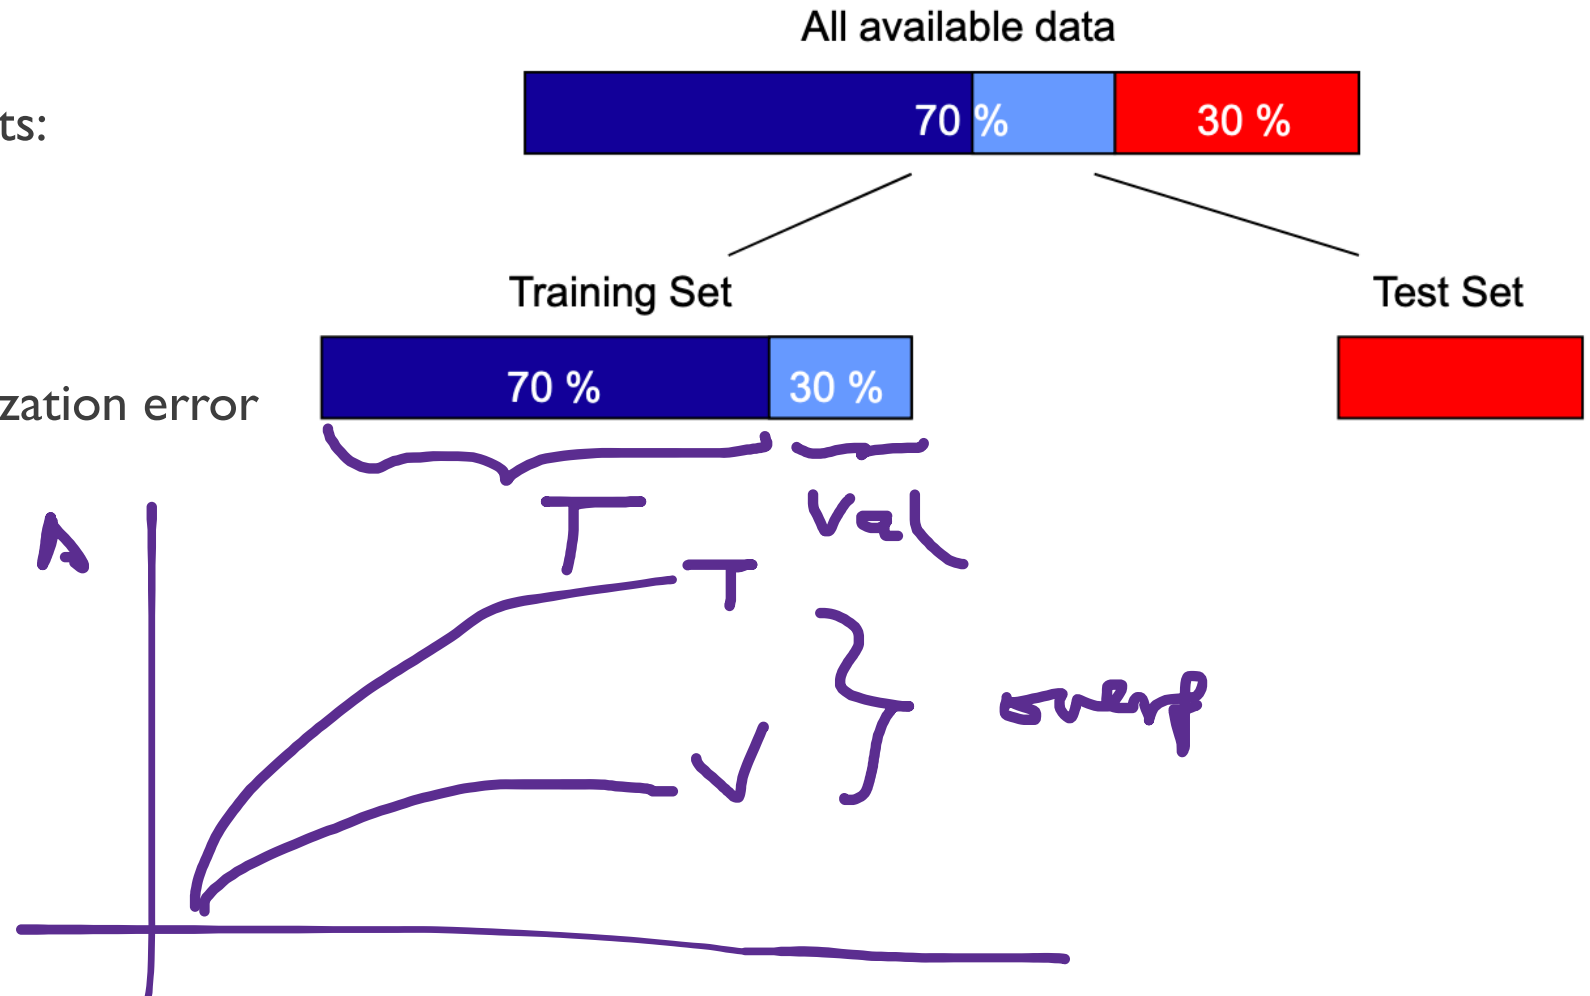
MODEL SELECTION

- Performed during model building
- Select a model that is not overly complex
 - (potential concerns for overly complex model: overfitting)
- Estimate generalization error
 - validation set
 - model complexity



MODEL SELECTION: USING VALIDATION SET

- Divide training data into two parts:
 - Training set:
 - Validation set:
 - use for estimating generalization error
- Drawback:
 - less data available for training



MODEL SELECTION: INCORPORATING MODEL COMPLEXITY

- Rationale: Occam's Razor
 - Given two models of similar generalization errors, one should prefer the simpler model
 - A complex model has a greater chance of overfitting
 - Include model complexity when evaluating a model


$$\text{Generalization Error}(\text{Model}) = \text{Train. Error}(\text{Model}, \text{Train. Data}) + \alpha \times \text{Complexity}(\text{Model})$$

ESTIMATING THE COMPLEXITY OF DECISION TREES


- **Pessimistic Error Estimate** of decision tree T with k leaf nodes:

$$err_{gen}(T) = \underbrace{err(T)} + \underbrace{\Omega}_{\text{Relative cost of adding a leaf node}} \times \frac{k}{N_{train}}$$

- $err(T)$: error rate on all training records
- Ω : trade-off hyper-parameter (similar to α)
 - Relative cost of adding a leaf node
- k : number of leaf nodes
- N_{train} : total number of training records




ESTIMATING THE COMPLEXITY OF DECISION TREES: EXAMPLE



Decision Tree, T_L

$$\begin{aligned} E_{p_L} &= \text{err.train} + I * k / N \\ &= 4/24 + 1 * 7 / 24 \end{aligned}$$



Decision Tree, T_R

$$\begin{aligned} E_{p_R} &= \text{error of train} + I * k / N \\ &= 6/24 + 1 * 4 / 24 \end{aligned}$$

$$err_{gen}(T) = err(T) + \Omega \times \frac{k}{N_{train}}$$

$$e(T_L) = 4/24$$

$$e(T_R) = 6/24$$

$$\Omega = 1$$


Pessimistic errors for both trees

$$e_{gen}(T_L) = 4/24 + 1 * 7/24 = 11/24 = 0.458$$


$$e_{gen}(T_R) = 6/24 + 1 * 4/24 = 10/24 = 0.417$$

ESTIMATING THE COMPLEXITY OF DECISION TREES

- Resubstitution Estimate:
 - **optimistic error** estimate: using training error as an estimate of generalization error



Decision Tree, T_1




Decision Tree, T_R

$$e(T_L) = 4/24$$

$$e(T_R) = 6/24$$

MINIMUM DESCRIPTION LENGTH (MDL)

X	y
X ₁	1
X ₂	0
X ₃	0
X ₄	1
...	...
X _n	1



X	y
X ₁	?
X ₂	?
X ₃	?
X ₄	?
...	...
X _n	?

- $\text{Cost}(\text{Model}, \text{Data}) = \text{Cost}(\text{Data}|\text{Model}) + \alpha \times \text{Cost}(\text{Model})$
 - Cost is the number of bits needed for encoding.
 - Search for the least costly model.
- $\text{Cost}(\text{Data}|\text{Model})$ encodes the misclassification errors.
- $\text{Cost}(\text{Model})$ uses node encoding (number of children) plus splitting condition encoding.

MODEL SELECTION FOR DECISION TREES

- **Pre-Pruning (Early Stopping Rule)**

- Stop the algorithm before it becomes a fully-grown tree

- Typical stopping conditions for a node:

- Stop if all instances belong to the same class

- or Stop if all the attribute values are the same

- More restrictive conditions:

- Stop if the number of instances is $<$ some user-specified threshold

- or Stop if class distribution of instances are independent of the available features (e.g., using χ^2 test)

- or Stop if expanding the current node does not improve impurity measures


- (e.g., Gini or information gain).

- or Stop if estimated generalization error falls below certain threshold

MODEL SELECTION FOR DECISION TREES

- **Post-pruning**

- Grow decision tree to its entirety
- Subtree replacement
 - Trim the nodes of the decision tree in a bottom-up fashion
 - If generalization error improves after trimming, replace sub-tree by a leaf node
 - Class label of leaf node is determined from **majority** class of instances in the sub-tree



EXAMPLE OF POST-PRUNING

Class = Yes	20
Class = No	10
Error = 10/30	

$$err_{gen}(T) = err(T) + \Omega \times \frac{k}{N_{train}}$$

Training Error (Before splitting) = 10/30

Pessimistic error = (10 + 0.5)/30 = 10.5/30

P_e = 10/30 + 1/8 * 4/30

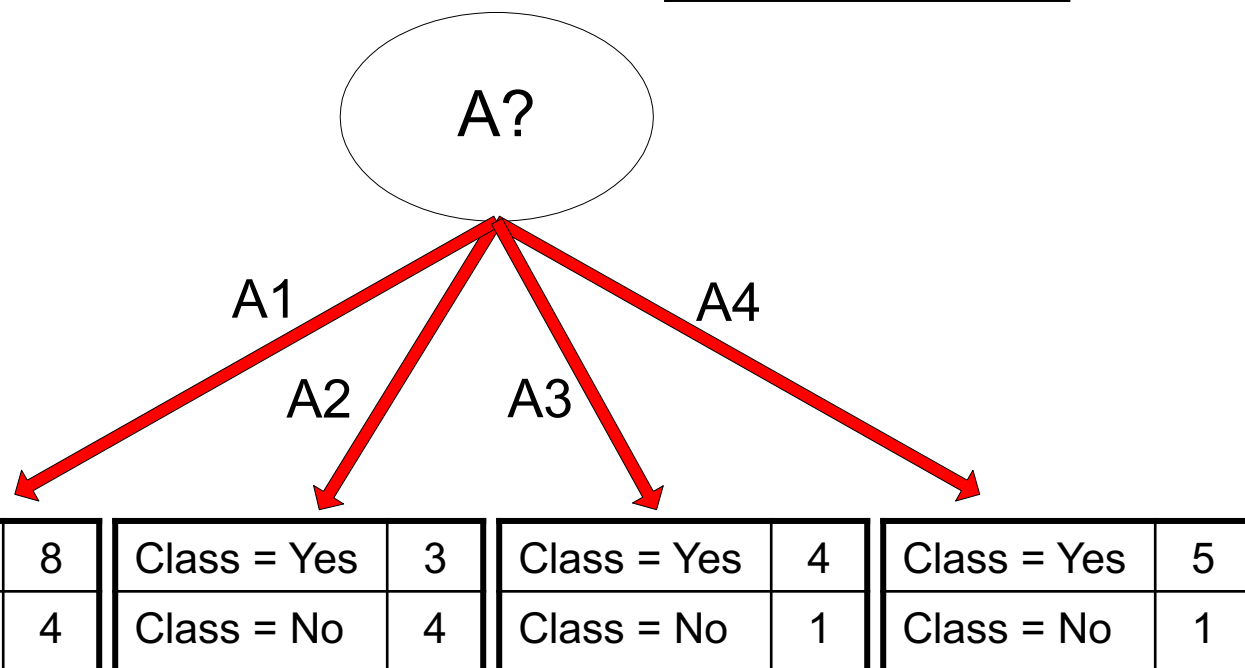
Training Error (After splitting) = 9/30

Pessimistic error (After splitting)


$$= (9 + 4 \times 0.5)/30 = \underline{11/30} > \frac{10.5}{30}$$

PRUNE!

$$GE(A) > GE(B)$$



EXAMPLE OF POST-PRUNING



EXAMPLES OF POST-PRUNING

Decision Tree:

```
depth = 1 :  
| breadth > 7 : class 1  
| breadth <= 7 :  
| | breadth <= 3 :  
| | | ImagePages > 0.375 : class 0  
| | | ImagePages <= 0.375 :  
| | | | totalPages <= 6 : class 1  
| | | | totalPages > 6 :  
| | | | | breadth <= 1 : class 1  
| | | | | breadth > 1 : class 0  
| | width > 3 :  
| | | MultiIP = 0:  
| | | | ImagePages <= 0.1333 : class 1  
| | | | ImagePages > 0.1333 :  
| | | | | breadth <= 6 : class 0  
| | | | | breadth > 6 : class 1  
| | | MultiIP = 1:  
| | | | TotalTime <= 361 : class 0  
| | | | TotalTime > 361 : class 1  
depth > 1 :  
| MultiAgent = 0:  
| | depth > 2 : class 0  
| | depth <= 2 :  
| | | MultiIP = 1: class 0  
| | | MultiIP = 0:  
| | | | breadth <= 6 : class 0  
| | | | breadth > 6 :  
| | | | | RepeatedAccess <= 0.0322 : class 0  
| | | | | RepeatedAccess > 0.0322 : class 1  
| MultiAgent = 1:
```

Subtree
Raising

Simplified Decision Tree:

```
depth = 1 :  
| | ImagePages <= 0.1333 : class 1  
| | ImagePages > 0.1333 :  
| | | breadth <= 6 : class 0  
| | | breadth > 6 : class 1  
depth > 1 :  
| MultiAgent = 0: class 0  
| MultiAgent = 1:  
| | totalPages <= 81 : class 0  
| | totalPages > 81 : class 1
```

Subtree
Replacement


MODEL EVALUATION

- Purpose: k -fold 3 -fold: 120 $\frac{40}{S_1}$ $\frac{40}{S_2}$ $\frac{40}{S_3}$
 - Estimate performance of classifier on test set
- Holdout 5 : 120 $\frac{24}{S_1}$ $\frac{24}{S_2}$ $\frac{24}{S_3}$ $\frac{24}{S_4}$ $\frac{24}{S_5}$
 - Reserve $k\%$ for training and $(100-k)\%$ for testing
 - Random subsampling: repeated holdout
- Cross validation
 - Partition data into k disjoint subsets
 - k -fold: train on $k-1$ partitions, test on the remaining one

id	data
1	+
2	+
	+
10	-

CROSS-VALIDATION EXAMPLE

- 3-fold cross-validation



VARIATIONS ON CROSS-VALIDATION

- Repeated cross-validation

- Perform cross-validation for multiple times

- Give an estimate of the variance of the generalization error

- Stratified cross-validation

- Guarantee the same percentage of class labels in training and test

- Good for imbalanced datasets and small samples

- Use nested cross-validation approach for model selection and evaluation